

# Breast

Oriol Planesas, Heribert Roig

10 de abril, 2022

```
library(keras)
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.1.2
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(pROC)
```

```
## Warning: package 'pROC' was built under R version 4.1.2
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

**1. Describe protein abundance and gene expression datasets. How many patients have data of both types available. Are there missing data from some of the datasets? Preprocess them if necessary.**

```
protein_abundance <- read.csv(params$file1, sep = ",")
gene <- read.csv(params$file2, sep = ",")
clinical <- read.csv(params$file3, sep = "\t")
copy <- read.csv(params$file4, sep = "\t")
```

Haremos un pequeño resumen de los primeros genes de las diversas bases de datos para ver en que consisten estos:

```
summary(gene[,1:10])
```

##	Sample	ELM02	CREB3L1	RPS11
##	Length:526	Min. :-1.52492	Min. :-1.4127	Min. :-1.02588
##	Class :character	1st Qu.: -0.04492	1st Qu.: 0.3743	1st Qu.: 0.08306
##	Mode :character	Median : 0.32621	Median : 0.9550	Median : 0.38512
##		Mean : 0.31785	Mean : 0.9259	Mean : 0.40055
##		3rd Qu.: 0.60642	3rd Qu.: 1.4921	3rd Qu.: 0.70337
##		Max. : 2.90117	Max. : 3.2203	Max. : 2.52200
##	PNMA1	MMP2	C10orf90	ZHX3
##	Min. :-1.72625	Min. :-2.2847	Min. :-5.027	Min. :-1.73750
##	1st Qu.: -0.38362	1st Qu.: -0.3141	1st Qu.: -2.233	1st Qu.: -0.48471
##	Median : 0.08175	Median : 0.4322	Median : -1.933	Median : -0.22008
##	Mean : 0.02557	Mean : 0.3678	Mean : -1.735	Mean : -0.21253

```
## 3rd Qu.: 0.45344 3rd Qu.: 1.0822 3rd Qu.: -1.542 3rd Qu.: 0.06704
## Max. : 1.60050 Max. : 2.8012 Max. : 2.572 Max. : 1.36167
## ERCC5 GPR98
## Min. : -1.51125 Min. : -2.8070
## 1st Qu.: -0.29169 1st Qu.: -1.4713
## Median : 0.02825 Median : -0.2177
## Mean : 0.03835 Mean : -0.4008
## 3rd Qu.: 0.38344 3rd Qu.: 0.4988
## Max. : 1.94175 Max. : 2.3814
```

```
summary(clinical[,1:4])
```

```
## Sample Histology PAM50Call
## Length:847 Length:847 Length:847
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
## ajcc_cancer_metastasis_stage_code
## Length:847
## Class :character
## Mode :character
```

```
summary(protein_abundance[,1:10])
```

```
## Sample X14.3.3_epsilon X4E.BP1 X4E.BP1_pS65
## Length:410 Min. : -0.46787 Min. : -1.19402 Min. : -0.845501
## Class :character 1st Qu.: -0.09601 1st Qu.: -0.28960 1st Qu.: -0.164652
## Mode :character Median : 0.01371 Median : -0.04515 Median : 0.005191
## Mean : 0.05439 Mean : 0.09174 Mean : 0.026804
## 3rd Qu.: 0.17245 3rd Qu.: 0.36732 3rd Qu.: 0.192625
## Max. : 0.93994 Max. : 2.56981 Max. : 1.123657
## X4E.BP1_pT37 X4E.BP1_pT70 X53BP1 A.Raf_pS299
## Min. : -1.2462 Min. : -0.72685 Min. : -1.72613 Min. : -1.0743520
## 1st Qu.: -0.3216 1st Qu.: -0.11489 1st Qu.: -0.30133 1st Qu.: -0.1954479
## Median : 0.1167 Median : 0.02684 Median : 0.01189 Median : -0.0146859
## Mean : 0.1404 Mean : 0.05617 Mean : 0.02492 Mean : 0.0000203
## 3rd Qu.: 0.4857 3rd Qu.: 0.18398 3rd Qu.: 0.37898 3rd Qu.: 0.1523013
## Max. : 2.6901 Max. : 1.25149 Max. : 2.65049 Max. : 1.0728245
## ACC1 ACC_pS79
## Min. : -2.45608 Min. : -1.59836
## 1st Qu.: -0.41713 1st Qu.: -0.35265
## Median : 0.06032 Median : 0.08966
## Mean : 0.11845 Mean : 0.17207
## 3rd Qu.: 0.60444 3rd Qu.: 0.63231
## Max. : 2.46339 Max. : 2.68644
```

Vemos como las bases de datos de protein y gene consisten en valores numericos con valores bajos alrededor del 0 tanto en positivo como negativo. También observamos como aquellos genes y proteínas con valores similares tienen también un nonmbre parecido, por consiguiente, parece que aquellos genes relacionados entre si o que son parecidos tienen un efecto similar.

```
set1<-intersect(protein_abundance$Sample,gene$Sample)
set1 <- intersect(set1, clinical$Sample)
```

```
xgene<-gene[gene$Sample %in% set1,]
xprotein<-protein_abundance[protein_abundance$Sample %in% set1,]
```

```
xclinical <- clinical[clinical$Sample %in% set1,]
xclinical <- xclinical[,c(1,9)]
```

```
dim(xgene)
```

```
## [1] 387 17815
```

387 individuos están presentes en los datasets clinical, gene y protein\_abundance.

```
sum(is.na(xgene)) # 1161 missings en gene
```

```
## [1] 1161
```

```
sum(is.na(xprotein)) # 0 missings en protein
```

```
## [1] 0
```

```
sum(is.na(xclinical)) # 0 missings en protein
```

```
## [1] 0
```

Podemos observar como en la base de datos “gene\_expression” hay 1161 valores missing, que arreglaremos mas adelante.

## With gene expression data:

### 2. Select the 25% of genes with the most variability

```
gene_var <- diag(var(gene[,2:ncol(gene)], na.rm = T)) # Calculamos la variabilidad de los genes
```

```
gene_topvar <- sort(gene_var, decreasing = T)[1:(length(gene_var)/20)]
```

```
# Cogemos solo el 5% con la mayor variabilidad para no tener problemas con los modelos
```

```
# Separamos los 5% de genes con mas variabilidad
```

```
xgene <- xgene[,c("Sample", names(gene_topvar))]
```

```
siNAcol <- apply(is.na(xgene), 2, sum) >= 1
```

```
xgene <- xgene[,!siNAcol]
```

```
sum(is.na(xgene)) # No hi ha missings
```

```
## [1] 0
```

```
dim(xgene)
```

```
## [1] 387 890
```

En lugar de utilizar el 25% de genes con más variabilidad, utilizaremos sólo el 5% debido a que el número de variables es relativamente grande para ejecutar el programa en nuestros ordenadores personales. Asimismo, obtenemos un conjunto de datos sin missings.

Los datos de gene expression que utilizaremos tienen una dimension final de 387 observaciones y 890 variables.

```
# Eliminamos las filas donde la respuesta no es ni negativa ni positiva:
```

```
sel1<-which(xclinical$breast_carcinoma_estrogen_receptor_status != "Positive")
```

```
sel2<-which(xclinical$breast_carcinoma_estrogen_receptor_status != "Negative")
```

```
sel<-intersect(sel1,sel2)
```

```
xclinical<-xclinical[-sel,]
data1 <- merge(xclinical, xgene, by.x = "Sample", by.y = "Sample")
```

Generamos los datos de training y test para los modelos que vengan a continuacion de la base de datos gene:

```
set.seed(123)
training<-sample(1:nrow(data1),2*nrow(data1)/3)

escalat1 <- scale(data1[,-c(1,2)])

xtrain<-escalat1[training,]
xtest<-escalat1[-training,]
ytrain<-data1[training,2]
ytest<-data1[-training,2]
ylabls<-vector()
ylabls[ytrain=="Positive"]<-1
ylabls[ytrain=="Negative"]<-0
ylabelstest<-vector()
ylabelstest[ytest=="Positive"]<-1
ylabelstest[ytest=="Negative"]<-0
```

### 3. Implement an stacked autoencoder (SAE) with three stacked layers of 1000, 100, 50 nodes. Provide in each case evidence of the quality of the coding obtained.

Empezaremos generando el primer autoencoder con un shape de el numero de columnas nummericas que hay en la base de datos xgene, que es el mismo numero de columnas que tiene el train, en el primer layer\_input con los datos de "gene\_expression". El primer decode tendra 1000 nodos y asi sucesivamente con los valores dados en el enunciado.

```
# Autoencoder 1
```

```
# Encoder
```

```
input_enc1 <- layer_input(shape = (ncol(xgene) - 1))
```

```
## Loaded Tensorflow version 2.6.0
```

```
output_enc1 <- input_enc1 %>%
  layer_dense(units = 1000, activation = "relu", name='G_Enc1')
encoder1 = keras_model(input_enc1, output_enc1)
summary(encoder1)
```

```
## Model: "model"
```

```
## -----
## Layer (type)                Output Shape          Param #
## -----
## input_1 (InputLayer)        [(None, 889)]         0
## -----
## G_Enc1 (Dense)              (None, 1000)          890000
## -----
## Total params: 890,000
## Trainable params: 890,000
## Non-trainable params: 0
## -----
```

```
# Decoder
```

```
input_dec1 = layer_input(shape = 1000)
output_dec1 <- input_dec1 %>%
```

```

layer_dense(units = (ncol(xgene)-1), activation="linear", name='G_Dec1')
decoder1 = keras_model(input_dec1, output_dec1)
summary(decoder1)

```

```

## Model: "model_1"
## -----
## Layer (type)                Output Shape          Param #
## =====
## input_2 (InputLayer)        [(None, 1000)]         0
## -----
## G_Dec1 (Dense)              (None, 889)           889889
## =====
## Total params: 889,889
## Trainable params: 889,889
## Non-trainable params: 0
## -----

```

```

# Juntar el encoder y el decoder
aen_input1 = layer_input(shape = (ncol(xgene)-1))
aen_output1 = aen_input1 %>%
  encoder1() %>%
  decoder1()
sae1 = keras_model(aen_input1, aen_output1)
summary(sae1)

```

```

## Model: "model_2"
## -----
## Layer (type)                Output Shape          Param #
## =====
## input_3 (InputLayer)        [(None, 889)]         0
## -----
## model (Functional)          (None, 1000)          890000
## -----
## model_1 (Functional)        (None, 889)           889889
## =====
## Total params: 1,779,889
## Trainable params: 1,779,889
## Non-trainable params: 0
## -----

```

```

sae1 %>% compile(
  optimizer = "rmsprop",
  loss = "mse")

sae1 %>% fit(
  x=as.matrix(xtrain),
  y=as.matrix(xtrain),
  epochs = 25,
  batch_size=64,
  validation_split = 0.2)

#Generador con en encoder
encoded_expression1 <- encoder1 %>% predict(as.matrix(xtrain))

```

El primer autoencoder tiene 1779889 parámetros.

```

# Autoencoder 2
input_enc2 <- layer_input(shape = 1000)
output_enc2 <- input_enc2 %>%
  layer_dense(units = 100, activation = "relu", name='Enc_AE2')
encoder2 = keras_model(input_enc2, output_enc2)
summary(encoder2)

## Model: "model_3"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_4 (InputLayer)        [(None, 1000)]              0
## -----
## Enc_AE2 (Dense)             (None, 100)                 100100
## =====
## Total params: 100,100
## Trainable params: 100,100
## Non-trainable params: 0
## -----

input_dec2 = layer_input(shape = 100)
output_dec2 <- input_dec2 %>%
  layer_dense(units = 1000, activation="linear", name='Dec_AE1')
decoder2 = keras_model(input_dec2, output_dec2)
summary(decoder2)

## Model: "model_4"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_5 (InputLayer)        [(None, 100)]              0
## -----
## Dec_AE1 (Dense)             (None, 1000)               101000
## =====
## Total params: 101,000
## Trainable params: 101,000
## Non-trainable params: 0
## -----

aen_input2 = input_enc2
aen_output2 = aen_input2 %>%
  encoder2() %>%
  decoder2()
sae2 = keras_model(aen_input2, aen_output2)
summary(sae2)

## Model: "model_5"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_4 (InputLayer)        [(None, 1000)]              0
## -----
## model_3 (Functional)        (None, 100)                 100100
## -----
## model_4 (Functional)        (None, 1000)               101000

```

```
## =====
## Total params: 201,100
## Trainable params: 201,100
## Non-trainable params: 0
## -----

sae2 %>% compile(
  optimizer = "rmsprop",
  loss = "mse")

sae2 %>% fit(
  x=as.matrix(encoded_expression1),
  y=as.matrix(encoded_expression1),
  epochs = 25,
  batch_size=64,
  validation_split = 0.2)

encoded_expression2 <- encoder2 %>% predict(as.matrix(encoded_expression1))
```

Este autoencoder tiene 201100 parámetros

```
# Autoencoder 3

# Encoder
input_enc3 <- layer_input(shape = 100)
output_enc3 <- input_enc3 %>%
  layer_dense(units = 50, activation = "relu", name='Enc_AE3')
encoder3 = keras_model(input_enc3, output_enc3)
summary(encoder3)
```

```
## Model: "model_6"
## -----
## Layer (type)                Output Shape                Param #
## -----
## input_6 (InputLayer)        [(None, 100)]              0
## -----
## Enc_AE3 (Dense)             (None, 50)                 5050
## -----
## Total params: 5,050
## Trainable params: 5,050
## Non-trainable params: 0
## -----
```

```
# Decoder
input_dec3 = layer_input(shape = 50)
output_dec3 <- input_dec3 %>%
  layer_dense(units = 100, activation="linear", name='Dec_AE1')
decoder3 = keras_model(input_dec3, output_dec3)
summary(decoder3)
```

```
## Model: "model_7"
## -----
## Layer (type)                Output Shape                Param #
## -----
## input_7 (InputLayer)        [(None, 50)]              0
## -----
```

```

## Dec_AE1 (Dense)                                (None, 100)                                5100
## =====
## Total params: 5,100
## Trainable params: 5,100
## Non-trainable params: 0
## -----
aen_input3 = input_enc3
aen_output3 = aen_input3 %>%
  encoder3() %>%
  decoder3()
sae3 = keras_model(aen_input3, aen_output3)
summary(sae3)

## Model: "model_8"
## -----
## Layer (type)                                Output Shape                                Param #
## =====
## input_6 (InputLayer)                        [(None, 100)]                              0
## -----
## model_6 (Functional)                        (None, 50)                                5050
## -----
## model_7 (Functional)                        (None, 100)                               5100
## =====
## Total params: 10,150
## Trainable params: 10,150
## Non-trainable params: 0
## -----

sae3 %>% compile(
  optimizer = "rmsprop",
  loss = "mse")

sae3 %>% fit(
  x=as.matrix(encoded_expression2),
  y=as.matrix(encoded_expression2),
  epochs = 40,
  batch_size=64,
  validation_split = 0.2)

encoded_expression3 <- encoder3 %>% predict(as.matrix(encoded_expression2))

```

El tercer autoencoder tiene 10150 parámetros.

**4.Using the SAE as pre-training model, couple it with a two-layer DNN to predict the state of the estrogen receptor. The DNN must have 10 nodes in the first layer followed by the output layer.**

Generamos el modelo juntando los 3 encoders anteriores

```

sae_input = layer_input(shape = (ncol(xgene)-1), name = "input_gene")
sae_output = sae_input %>%
  encoder1() %>%
  encoder2() %>%
  encoder3() %>%
  layer_dense(10, activation = "relu", name='L1_SAE1') %>%
  layer_dense(1, activation = "sigmoid", name='L2_SAE1')

```



```
sae = keras_model(sae_input, sae_output)
summary(sae)
```

```
## Model: "model_9"
## -----
## Layer (type)                Output Shape          Param #
## =====
## input_gene (InputLayer)      [(None, 889)]         0
## -----
## model (Functional)           (None, 1000)          890000
## -----
## model_3 (Functional)         (None, 100)           100100
## -----
## model_6 (Functional)         (None, 50)            5050
## -----
## L1_SAE1 (Dense)              (None, 10)            510
## -----
## L2_SAE1 (Dense)              (None, 1)             11
## =====
## Total params: 995,671
## Trainable params: 995,671
## Non-trainable params: 0
## -----
```

```
freeze_weights(sae,from=1,to=3)
```

Al juntar todos los encoder en un mismo modelo tenemos que este acaba teniendo 995671 parametros.

```
sae %>% compile(
  optimizer = "rmsprop",
  loss = 'binary_crossentropy',
  metric = "acc"
)
```

```
sae %>% fit(
  x=xtrain,
  y=ylabels,
  epochs = 15,
  batch_size = 64,
  validation_split = 0.2
)
```

El valor de accuracy que obtenemos es cercano a 0.85 y la pérdida es aproximadamente 0.4

```
sae %>% evaluate(as.matrix(xtest), ylabelstest)
```

```
##      loss      acc
## 0.3277466 0.8897638
```

Cuando evaluamos el modelo con los datos de test, conseguimos unos valores de las métricas de loss y accuracy muy parecidos a los de entrenamiento, por lo que podemos decir que tenemos un buen rendimiento en el modelo.

```
yhat <- predict(sae,as.matrix(xtest))
```

```
yhatclass<-as.factor(ifelse(yhat<0.5,0,1))
table(yhatclass, ylabelstest)
```

```
##          ylabelstest
## yhatclass 0  1
##          0 20  6
##          1  8 93

confusionMatrix(yhatclass,as.factor(ylabelstest))
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction 0  1
##          0 20  6
##          1  8 93
##
##          Accuracy : 0.8898
##          95% CI : (0.822, 0.9384)
##    No Information Rate : 0.7795
##    P-Value [Acc > NIR] : 0.001011
##
##          Kappa : 0.6709
##
## Mcnemar's Test P-Value : 0.789268
##
##          Sensitivity : 0.7143
##          Specificity : 0.9394
##          Pos Pred Value : 0.7692
##          Neg Pred Value : 0.9208
##          Prevalence : 0.2205
##          Detection Rate : 0.1575
##    Detection Prevalence : 0.2047
##          Balanced Accuracy : 0.8268
##
##          'Positive' Class : 0
##
```

Vemos que al predecir valores con la predicción observamos como el modelo tiene mayor error a la hora de predecir los casos negativos (0). Puede ser debido a que hay un número mayor de muestras con respuesta positiva, por lo que el modelo está más entrenado para este caso.

### 5. On the test set, provide the ROC curve and AUC and other performance metrics.


```
roc_sae_test <- roc(response = ylabelstest, predictor = yhat)

## Setting levels: control = 0, case = 1

## Warning in roc.default(response = ylabelstest, predictor = yhat): Deprecated use
## a matrix as predictor. Unexpected results may be produced, please pass a numeric
## vector.

## Setting direction: controls < cases

plot(roc_sae_test, col = "blue", print.auc=TRUE)
legend("bottomright", legend = c("sae"), lty = c(1), col = c("blue"))
```



En este grafico para ver el valor de auc obtenemos que este valor es de 0.915, teniendo asi que este modelo tiene un buen valor de diagnostico.

**6. With `tfruns()` repeat points 4 and 5, exploring the configurations of the first layer of the DNN based on 5, 10 and 20 nodes. Determine which configuration is the best.**

Para realizar el `tfruns`, generaremos el código en otro archivo `.R` y entonces cargaremos aqui los diferentes modelos con el codigo que hay a continuacion.

```
library(tfruns)
```

```
## Warning: package 'tfruns' was built under R version 4.1.2
```

```
nodes <- c(5, 10, 20)
```

```
for (i in 1:length(nodes)){
  print(i)
  training_run("Breast_tfruns.R",
              flags = c(units = nodes[i]))
}
```

```
## [1] 1
```

```
## Using run directory runs/2022-04-10T19-30-26Z
```

```
##
```

```
## > FLAGS <- flags(flag_integer("units", 10))
```

```
##
```

```
## > sae_input = layer_input(shape = (ncol(xgene) - 1))
```

```
##
```

```

## > sae_output = sae_input %>% encoder1() %>% encoder2() %>%
## +   encoder3() %>% layer_dense(FLAGS$units, activation = "relu") %>%
## +   layer_dense( .... [TRUNCATED]
##
## > sae = keras_model(sae_input, sae_output)
##
## > summary(sae)
## Model: "model_10"
##
## -----
## Layer (type)                Output Shape          Param #
## -----
## input_8 (InputLayer)        [(None, 889)]         0
## -----
## model (Functional)           (None, 1000)          890000
## -----
## model_3 (Functional)         (None, 100)           100100
## -----
## model_6 (Functional)         (None, 50)            5050
## -----
## dense_1 (Dense)              (None, 5)             255
## -----
## dense (Dense)                (None, 1)             6
## -----
## Total params: 995,411
## Trainable params: 5,311
## Non-trainable params: 990,100
## -----
##
## > sae %>% compile(optimizer = "rmsprop", loss = "binary_crossentropy",
## +   metric = "acc")
##
## > sae %>% fit(x = xtrain, y = ylabelstest, epochs = 15,
## +   batch_size = 64, validation_split = 0.2)
##
## > yhat <- predict(sae, as.matrix(xtest))
##
## > yhatclass <- as.factor(ifelse(yhat < 0.5, 0, 1))
##
## > table(yhatclass, ylabelstest)
##      ylabelstest
## yhatclass 0  1
##      0 22  7
##      1  6 92
##
## > confusionMatrix(yhatclass, as.factor(ylabelstest))
## Confusion Matrix and Statistics
##
##           Reference
## Prediction 0  1
##      0 22  7
##      1  6 92
##
##           Accuracy : 0.8976
##           95% CI : (0.8313, 0.9444)

```

```

##      No Information Rate : 0.7795
##      P-Value [Acc > NIR] : 0.0004164
##
##              Kappa : 0.706
##
## Mcnemar's Test P-Value : 1.0000000
##
##      Sensitivity : 0.7857
##      Specificity : 0.9293
##      Pos Pred Value : 0.7586
##      Neg Pred Value : 0.9388
##      Prevalence : 0.2205
##      Detection Rate : 0.1732
##      Detection Prevalence : 0.2283
##      Balanced Accuracy : 0.8575
##
##      'Positive' Class : 0
##
##
## > roc_sae_test <- roc(response = ylabelstest, predictor = yhat)
## Setting levels: control = 0, case = 1
## Warning in roc.default(response = ylabelstest, predictor = yhat): Deprecated use
## a matrix as predictor. Unexpected results may be produced, please pass a numeric
## vector.
## Setting direction: controls < cases
##
## > plot(roc_sae_test, col = "blue", print.auc = TRUE)
##
## > legend("bottomright", legend = c("sae"), lty = c(1),
## +      col = c("blue"))
##
## Run completed: runs/2022-04-10T19-30-26Z
## [1] 2
## Using run directory runs/2022-04-10T19-30-30Z
##
## > FLAGS <- flags(flag_integer("units", 10))
##
## > sae_input = layer_input(shape = (ncol(xgene) - 1))
##
## > sae_output = sae_input %>% encoder1() %>% encoder2() %>%
## +      encoder3() %>% layer_dense(FLAGS$units, activation = "relu") %>%
## +      layer_dense( ... [TRUNCATED]
##
## > sae = keras_model(sae_input, sae_output)
##
## > summary(sae)
## Model: "model"
## -----
## Layer (type)                                Output Shape                                Param #

```

```

## =====
## input_1 (InputLayer)          [(None, 889)]          0
## -----
## model (Functional)            (None, 1000)          890000
## -----
## model_3 (Functional)          (None, 100)          100100
## -----
## model_6 (Functional)          (None, 50)           5050
## -----
## dense_1 (Dense)              (None, 10)           510
## -----
## dense (Dense)                (None, 1)            11
## =====
## Total params: 995,671
## Trainable params: 5,571
## Non-trainable params: 990,100
## -----
##
## > sae %>% compile(optimizer = "rmsprop", loss = "binary_crossentropy",
## +   metric = "acc")
##
## > sae %>% fit(x = xtrain, y = ylabels, epochs = 15,
## +   batch_size = 64, validation_split = 0.2)
##
## > yhat <- predict(sae, as.matrix(xtest))
##
## > yhatclass <- as.factor(ifelse(yhat < 0.5, 0, 1))
##
## > table(yhatclass, ylabelstest)
##      ylabelstest
## yhatclass 0  1
##      0 21  6
##      1  7 93
##
## > confusionMatrix(yhatclass, as.factor(ylabelstest))
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 21  6
##           1  7 93
##
##              Accuracy : 0.8976
##              95% CI : (0.8313, 0.9444)
##      No Information Rate : 0.7795
##      P-Value [Acc > NIR] : 0.0004164
##
##              Kappa : 0.6983
##
## Mcnemar's Test P-Value : 1.0000000
##
##              Sensitivity : 0.7500
##              Specificity : 0.9394
##              Pos Pred Value : 0.7778

```

```

##          Neg Pred Value : 0.9300
##          Prevalence : 0.2205
##          Detection Rate : 0.1654
##          Detection Prevalence : 0.2126
##          Balanced Accuracy : 0.8447
##
##          'Positive' Class : 0
##
##
## > roc_sae_test <- roc(response = ylabelstest, predictor = yhat)
## Setting levels: control = 0, case = 1
## Warning in roc.default(response = ylabelstest, predictor = yhat): Deprecated use
## a matrix as predictor. Unexpected results may be produced, please pass a numeric
## vector.
## Setting direction: controls < cases
##
## > plot(roc_sae_test, col = "blue", print.auc = TRUE)
##
## > legend("bottomright", legend = c("sae"), lty = c(1),
## +      col = c("blue"))
##
## Run completed: runs/2022-04-10T19-30-30Z
## [1] 3
## Using run directory runs/2022-04-10T19-30-34Z
##
## > FLAGS <- flags(flag_integer("units", 10))
##
## > sae_input = layer_input(shape = (ncol(xgene) - 1))
##
## > sae_output = sae_input %>% encoder1() %>% encoder2() %>%
## +   encoder3() %>% layer_dense(FLAGS$units, activation = "relu") %>%
## +   layer_dense( ... [TRUNCATED]
##
## > sae = keras_model(sae_input, sae_output)
##
## > summary(sae)
## Model: "model"
## -----
## Layer (type)                Output Shape                Param #
## -----
## input_1 (InputLayer)        [(None, 889)]              0
## -----
## model (Functional)           (None, 1000)               890000
## -----
## model_3 (Functional)         (None, 100)                100100
## -----
## model_6 (Functional)         (None, 50)                 5050
## -----
## dense_1 (Dense)              (None, 20)                 1020

```

```

## -----
## dense (Dense) (None, 1) 21
## =====
## Total params: 996,191
## Trainable params: 6,091
## Non-trainable params: 990,100
## -----
##
## > sae %>% compile(optimizer = "rmsprop", loss = "binary_crossentropy",
## +   metric = "acc")
##
## > sae %>% fit(x = xtrain, y = ylabels, epochs = 15,
## +   batch_size = 64, validation_split = 0.2)
##
## > yhat <- predict(sae, as.matrix(xtest))
##
## > yhatclass <- as.factor(ifelse(yhat < 0.5, 0, 1))
##
## > table(yhatclass, ylabelstest)
##      ylabelstest
## yhatclass 0 1
##      0 22 4
##      1 6 95
##
## > confusionMatrix(yhatclass, as.factor(ylabelstest))
## Confusion Matrix and Statistics
##
##      Reference
## Prediction 0 1
##      0 22 4
##      1 6 95
##
##      Accuracy : 0.9213
##      95% CI : (0.86, 0.9616)
##      No Information Rate : 0.7795
##      P-Value [Acc > NIR] : 1.759e-05
##
##      Kappa : 0.7649
##
##      McNemar's Test P-Value : 0.7518
##
##      Sensitivity : 0.7857
##      Specificity : 0.9596
##      Pos Pred Value : 0.8462
##      Neg Pred Value : 0.9406
##      Prevalence : 0.2205
##      Detection Rate : 0.1732
##      Detection Prevalence : 0.2047
##      Balanced Accuracy : 0.8727
##
##      'Positive' Class : 0
##
## > roc_sae_test <- roc(response = ylabelstest, predictor = yhat)

```



```
## Setting levels: control = 0, case = 1

## Warning in roc.default(response = ylabelstest, predictor = yhat): Deprecated use
## a matrix as predictor. Unexpected results may be produced, please pass a numeric
## vector.

## Setting direction: controls < cases

##
## > plot(roc_sae_test, col = "blue", print.auc = TRUE)

##
## > legend("bottomright", legend = c("sae"), lty = c(1),
## +      col = c("blue"))

##
## Run completed: runs/2022-04-10T19-30-34Z

runs <- ls_runs(latest_n = 3)
runs <- runs[, c("flag_units", "metric_val_acc", "metric_val_loss")]
(runs <- runs[order(runs$metric_val_acc, decreasing = T),])
```

```
## Data frame: 3 x 3
##   flag_units metric_val_acc metric_val_loss
## 1         20         0.8824         0.4612
## 3          5         0.8824         0.4392
## 2         10         0.8627         0.5518
```

Vemos que los tres modelos ejecutados tienen un accuracy parecido, por lo que en los siguientes ejercicios utilizaremos la configuración inicial.

So far, we have two SAEs. One for the abundance of proteins (see class examples) and the other for gene expression we just built.

**7. Split the set of patients with complete data (gene expression and protein abundance) in train and test sets.**

```
xprotein<-protein_abundance[protein_abundance$Sample %in% set1,]

data2 <- merge(xclinical, xprotein, by.x = "Sample", by.y = "Sample")
data2 <- merge(data2, xgene, by.x = "Sample", by.y = "Sample")

escalat2 <- scale(data2[, -c(1,2)])

xtrain2<-escalat2[training, -c(1,2)]
xtest2<-escalat2[-training, -c(1,2)]
xtrain2<-scale(xtrain2)
xtest2<-scale(xtest2)
ytrain2<-escalat2[training, 2]
ytest2<-escalat2[-training, 2]
ylabelstest2<-vector()
ylabelstest2[ytrain2=="Positive"]<-1
ylabelstest2[ytrain2=="Negative"]<-0
ylabelstest2<-vector()
ylabelstest2[ytest2=="Positive"]<-1
ylabelstest2[ytest2=="Negative"]<-0
```

**8. Concatenate the two SAEs to fit, on the trainset, a DNN that integrates both data sources to predict estrogen receptor status. The DNN must have a dense layer (with the better number**

of nodes according with point 6) and the output layer.

Modelo de la proteína

```
data3<-merge(xclinical,xprotein,by.x="Sample",by.y="Sample")

escalat3 <- scale(data3[, -c(1,2)])

xtrain3<-escalat3[training,]
xtest3<-escalat3[-training,]

ytrain3<-data3[training,2]
ytest3<-data3[-training,2]

ylabels3<-vector()
ylabels3[ytrain3=="Positive"]<-1
ylabels3[ytrain3=="Negative"]<-0

ytestlabels3<-vector()
ytestlabels3[ytest3=="Positive"]<-1
ytestlabels3[ytest3=="Negative"]<-0

# AE1
input_enc1_prot<-layer_input(shape = 142)
output_enc1_prot<-input_enc1_prot %>%
  layer_dense(units=50,activation="relu")
encoder1_prot = keras_model(input_enc1_prot, output_enc1_prot, name = "AE1")
summary(encoder1_prot)

## Model: "AE1"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_1 (InputLayer)        [(None, 142)]               0
## -----
## dense (Dense)               (None, 50)                  7150
## =====
## Total params: 7,150
## Trainable params: 7,150
## Non-trainable params: 0
## -----

input_dec1_prot = layer_input(shape = 50)
output_dec1_prot<-input_dec1_prot %>%
  layer_dense(units=142,activation="linear")

decoder1_prot = keras_model(input_dec1_prot, output_dec1_prot)

summary(decoder1_prot)

## Model: "model"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_2 (InputLayer)        [(None, 50)]                0
```

```

## -----
## dense_1 (Dense)                                (None, 142)                                7242
## =====
## Total params: 7,242
## Trainable params: 7,242
## Non-trainable params: 0
## -----

aen_input1_prot = layer_input(shape = 142)
aen_output1_prot = aen_input1_prot %>%
  encoder1_prot() %>%
  decoder1_prot()

sae1_prot = keras_model(aen_input1_prot, aen_output1_prot)
summary(sae1_prot)

## Model: "model_1"
## -----
## Layer (type)                                Output Shape                                Param #
## =====
## input_3 (InputLayer)                        [(None, 142)]                                0
## -----
## AE1 (Functional)                            (None, 50)                                7150
## -----
## model (Functional)                          (None, 142)                                7242
## =====
## Total params: 14,392
## Trainable params: 14,392
## Non-trainable params: 0
## -----

sae1_prot %>% compile(
  optimizer = "rmsprop",
  loss = "mse"
)

sae1_prot %>% fit(
  x=as.matrix(xtrain3),
  y=as.matrix(xtrain3),
  epochs = 50,
  batch_size=64,
  validation_split = 0.2
)

#Generating with Autoencoder
encoded_expression1_prot <- encoder1_prot %>% predict(as.matrix(xtrain3))

```

El primer autoencoder tiene 14392 parámetros

```

# AE2
input_enc2_prot<-layer_input(shape = 50)
output_enc2_prot<-input_enc2_prot %>%
  layer_dense(units=20,activation="relu")
encoder2_prot = keras_model(input_enc2_prot, output_enc2_prot)
summary(encoder2_prot)

```

```
## Model: "model_2"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_4 (InputLayer)         [(None, 50)]                0
## -----
## dense_2 (Dense)              (None, 20)                  1020
## =====
## Total params: 1,020
## Trainable params: 1,020
## Non-trainable params: 0
## -----
```

```
input_dec2_prot = layer_input(shape = 20)
output_dec2_prot<-input_dec2_prot %>%
  layer_dense(units=50,activation="linear")

decoder2_prot = keras_model(input_dec2_prot, output_dec2_prot)

summary(decoder2_prot)
```

```
## Model: "model_3"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_5 (InputLayer)         [(None, 20)]                0
## -----
## dense_3 (Dense)              (None, 50)                  1050
## =====
## Total params: 1,050
## Trainable params: 1,050
## Non-trainable params: 0
## -----
```

```
aen_input2_prot = layer_input(shape = 50)
aen_output2_prot = aen_input2_prot %>%
  encoder2_prot() %>%
  decoder2_prot()

sae2_prot = keras_model(aen_input2_prot, aen_output2_prot)
summary(sae2_prot)
```

```
## Model: "model_4"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_6 (InputLayer)         [(None, 50)]                0
## -----
## model_2 (Functional)          (None, 20)                  1020
## -----
## model_3 (Functional)          (None, 50)                  1050
## =====
## Total params: 2,070
## Trainable params: 2,070
## Non-trainable params: 0
```

```
## -----
sae2_prot %>% compile(
  optimizer = "rmsprop",
  loss = "mse"
)

sae2_prot %>% fit(
  x=as.matrix(encoded_expression1_prot),
  y=as.matrix(encoded_expression1_prot),
  epochs = 50,
  batch_size=64,
  validation_split = 0.2
)

#Generating with Autoencoder
encoded_expression2_prot <- encoder2_prot %>% predict(as.matrix(encoded_expression1_prot))
```

El segundo autoencoder tiene 2070 parámetros.

```
# AE3
input_enc3_prot<-layer_input(shape = 20)
output_enc3_prot<-input_enc3_prot %>%
  layer_dense(units=10,activation="relu")
encoder3_prot = keras_model(input_enc3_prot, output_enc3_prot)
summary(encoder3_prot)
```

```
## Model: "model_5"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_7 (InputLayer)        [(None, 20)]                0
## -----
## dense_4 (Dense)             (None, 10)                  210
## =====
## Total params: 210
## Trainable params: 210
## Non-trainable params: 0
## -----
```

```
input_dec3_prot = layer_input(shape = 10)
output_dec3_prot<-input_dec3_prot %>%
  layer_dense(units=20,activation="linear")

decoder3_prot = keras_model(input_dec3_prot, output_dec3_prot)

summary(decoder3_prot)
```

```
## Model: "model_6"
## -----
## Layer (type)                Output Shape                Param #
## =====
## input_8 (InputLayer)        [(None, 10)]                0
## -----
## dense_5 (Dense)             (None, 20)                  220
## =====
```

```

## Total params: 220
## Trainable params: 220
## Non-trainable params: 0
## -----
aen_input3_prot = layer_input(shape = 20)
aen_output3_prot = aen_input3_prot %>%
  encoder3_prot() %>%
  decoder3_prot()

sae3_prot = keras_model(aen_input3_prot, aen_output3_prot)
summary(sae3_prot)

## Model: "model_7"
## -----
## Layer (type)                Output Shape          Param #
## =====
## input_9 (InputLayer)        [(None, 20)]          0
## -----
## model_5 (Functional)        (None, 10)            210
## -----
## model_6 (Functional)        (None, 20)            220
## =====
## Total params: 430
## Trainable params: 430
## Non-trainable params: 0
## -----
sae3_prot %>% compile(
  optimizer = "rmsprop",
  loss = "mse"
)

sae3_prot %>% fit(
  x=as.matrix(encoded_expression2_prot),
  y=as.matrix(encoded_expression2_prot),
  epochs = 50,
  batch_size=64,
  validation_split = 0.2
)

#Generating with Autoencoder
encoded_expression3_prot <- encoder3_prot %>% predict(as.matrix(encoded_expression2_prot))

```

El tercer autoencoder para el conjunto de datos proteicos tiene 430 parámetros.

```

### Final model

sae_input_prot = layer_input(shape = 142, name = "input_prot")
sae_output_prot = sae_input_prot %>%
  encoder1_prot() %>%
  encoder2_prot() %>%
  encoder3_prot() %>%
  layer_dense(5, activation = "relu") %>%
  layer_dense(1, activation = "sigmoid")

```

```
sae_prot = keras_model(sae_input_prot, sae_output_prot)
summary(sae_prot)
```

```
## Model: "model_8"
## -----
## Layer (type)                Output Shape          Param #
## =====
## input_prot (InputLayer)      [(None, 142)]         0
## -----
## AE1 (Functional)             (None, 50)            7150
## -----
## model_2 (Functional)         (None, 20)            1020
## -----
## model_5 (Functional)         (None, 10)            210
## -----
## dense_7 (Dense)              (None, 5)             55
## -----
## dense_6 (Dense)              (None, 1)             6
## =====
## Total params: 8,441
## Trainable params: 8,441
## Non-trainable params: 0
## -----
```

El total de parámetros para el modelo Stacked autoencoder para el conjunto de datos de protein\_abundance es de 8170.

```
freeze_weights(sae_prot,from=1,to=3)
```

```
sae_prot %>% compile(
  optimizer = "rmsprop",
  loss = 'binary_crossentropy',
  metric = "acc"
)
```

```
sae_prot %>% fit(
  x=xtrain3,
  y=ylabels3,
  epochs = 30,
  batch_size=64,
  validation_split = 0.2
)
```

```
sae_prot %>%
  evaluate(as.matrix(xtest3), ytestlabels3)
```

```
##      loss      acc
## 0.3472343 0.7795275
```

Para este modelo, la precisión en la evaluación del conjunto de test es cercana a 0.90.

```
yhat_prot <- predict(sae_prot,as.matrix(xtest3))
```

```
yhatclass_prot<-as.factor(ifelse(yhat_prot<0.5,0,1))
table(yhatclass_prot, ytestlabels3)
```

```
##                ytestlabels3
```

```
## yhatclass_prot 0 1
##               1 28 99
```

Vemos que el porcentaje de error es parecido en ambos casos.

```
confusionMatrix(yhatclass_prot,as.factor(ytestlabels3))
```

```
## Warning in confusionMatrix.default(yhatclass_prot, as.factor(ytestlabels3)):
## Levels are not in the same order for reference and data. Refactoring data to
## match.
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction 0 1
```

```
##           0 0 0
```

```
##           1 28 99
```

```
##
```

```
##           Accuracy : 0.7795
```

```
##           95% CI : (0.6974, 0.8482)
```

```
## No Information Rate : 0.7795
```

```
## P-Value [Acc > NIR] : 0.5504
```

```
##
```

```
##           Kappa : 0
```

```
##
```

```
## McNemar's Test P-Value : 3.352e-07
```

```
##
```

```
##           Sensitivity : 0.0000
```

```
##           Specificity : 1.0000
```

```
## Pos Pred Value : NaN
```

```
## Neg Pred Value : 0.7795
```

```
## Prevalence : 0.2205
```

```
## Detection Rate : 0.0000
```

```
## Detection Prevalence : 0.0000
```

```
## Balanced Accuracy : 0.5000
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

Concatenate the 2 models:

```
sae_input_prova1 = layer_input(shape = (ncol(xgene)-1), name = "input_gene_prova")
```

```
sae_output_prova1 = sae_input_prova1 %>%
```

```
  encoder1() %>%
```

```
  encoder2() %>%
```

```
  encoder3() %>%
```

```
  layer_dense(10,activation = "relu", name='L1_SAE1')%>%
```

```
  layer_dense(1,activation = "sigmoid", name='L2_SAE1')
```

```
sae_prova1 = keras_model(sae_input_prova1, sae_output_prova1)
```

```
sae_input_prova2 = layer_input(shape = 142, name = "input_prot_prova")
```

```
sae_output_prova2 = sae_input_prova2 %>%
```

```
  encoder1_prot() %>%
```

```
  encoder2_prot() %>%
```

```
  encoder3_prot() %>%
```

```
  layer_dense(10,activation = "relu", name='L1_SAE2')%>%
```

```
  layer_dense(1,activation = "sigmoid", name='L2_SAE2')
```



```
sae_prova2 = keras_model(sae_input_prova2, sae_output_prova2)

concatenated<-layer_concatenate(list(sae_output_prova1,sae_output_prova2))

model_output_con<-concatenated %>%
  layer_dense(units = 20,"relu") %>%
  layer_dense(units = 1,activation = "sigmoid")

model_final<-keras_model(list(sae_input_prova1, sae_input_prova2), model_output_con)
summary(model_final)
```

```
## Model: "model_11"
##
## -----
## Layer (type)           Output Shape      Param #   Connected to
## -----
## input_gene_prova (InputLa [(None, 889)]      0
##
## input_prot_prova (InputLa [(None, 142)]      0
##
## model (Functional)      (None, 1000)      890000    input_gene_prova[0][0]
##
## AE1 (Functional)        (None, 50)        7150      input_prot_prova[0][0]
##
## model_3 (Functional)    (None, 100)       100100     model[5][0]
##
## model_2 (Functional)    (None, 20)        1020      AE1[2][0]
##
## model_6 (Functional)    (None, 50)        5050      model_3[5][0]
##
## model_5 (Functional)    (None, 10)        210       model_2[2][0]
##
## L1_SAE1 (Dense)        (None, 10)        510       model_6[5][0]
##
## L1_SAE2 (Dense)        (None, 10)        110       model_5[2][0]
##
## L2_SAE1 (Dense)        (None, 1)         11        L1_SAE1[0][0]
##
## L2_SAE2 (Dense)        (None, 1)         11        L1_SAE2[0][0]
##
## concatenate (Concatenate) (None, 2)         0         L2_SAE1[0][0]
##                                     L2_SAE2[0][0]
##
## dense_9 (Dense)        (None, 20)        60        concatenate[0][0]
##
## dense_8 (Dense)        (None, 1)         21        dense_9[0][0]
##
## =====
## Total params: 1,004,253
## Trainable params: 5,983
## Non-trainable params: 998,270
## -----
```

En el modelo concatenado tenemos un total de 1004253 parámetros.

```

model_final %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = "acc"
)

# training

model_final %>% fit(
  x = list(input_gene_prova = as.matrix(xtrain), input_prot_prova = as.matrix(xtrain3)),
  y = array(ylabels), epochs = 30, batch_size = 64, validation_split = 0.2
)

```

```

model_final %>%
  evaluate(list(as.matrix(xtest), as.matrix(xtest3)), ylabelstest)

```

```

##      loss      acc
## 0.4204175 0.9448819

```

En el modelo concatenado, el valor de pérdida es aproximadamente 0.50 y el de accuracy superior a 0.85.

```

yhat_final <- predict(model_final,list(as.matrix(xtest), as.matrix(xtest3)))

```

```

yhatclass_final<-as.factor(ifelse(yhat_final<0.5,0,1))
table(yhatclass_final, ylabelstest)

```

```

##           ylabelstest
## yhatclass_final  0  1
##                0 25  4
##                1  3 95

```

```

confusionMatrix(yhatclass_final,as.factor(ylabelstest))

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 25  4
##           1  3 95
##
##              Accuracy : 0.9449
##              95% CI : (0.8897, 0.9776)
##      No Information Rate : 0.7795
##      P-Value [Acc > NIR] : 2.95e-07
##
##              Kappa : 0.8417
##
##  Mcnemar's Test P-Value : 1
##
##              Sensitivity : 0.8929
##              Specificity : 0.9596
##      Pos Pred Value : 0.8621
##      Neg Pred Value : 0.9694
##              Prevalence : 0.2205
##      Detection Rate : 0.1969

```

```
## Detection Prevalence : 0.2283
## Balanced Accuracy : 0.9262
##
## 'Positive' Class : 0
##
```

9. On the testset, provide the ROC curve and AUC, and compare it with the model found in point 5.


```
roc_sae_test2 <- roc(response = ylabelstest, predictor = as.vector(yhat_final))
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_sae_test2, col = "blue", print.auc=TRUE)
```

```
legend("bottomright", legend = c("sae"), lty = c(1), col = c("blue"))
```



Segun este modelo y el valor obtenido de auc: 0.918, obtenemos que este modelo tiene una precisión alta para nuevos valores.


```
par(mfrow = c(1,2))
```

```
plot(roc_sae_test, col = "blue", print.auc=TRUE)
```

```
legend("bottomright", legend = c("sae"), lty = c(1), col = c("blue"))
```

```
plot(roc_sae_test2, col = "blue", print.auc=TRUE)
```

```
legend("bottomright", legend = c("sae"), lty = c(1), col = c("blue"))
```



```
par(mfrow = c(1,1))
```

Vemos que en los dos modelos obtenemos valores similares de AUC.

## 10. Discuss the results of the analysis.

Primero de todo tenemos el modelo del apartado 4, donde después de generar los 3 autoencoders tenemos 995671 parámetros. En este modelo, al compilarlo y entrenarlo con 15 “epochs” y un “batch size” de 64, obtenemos que al evaluarlo, el valor de la pérdida es de 0.40 y la precisión de 0.85. Con esto, podríamos decir que este modelo predice bastante bien.

En el apartado 6 comparamos diferentes capas y vemos como la diferencia entre los modelos es muy baja.

Finalmente, tenemos el modelo combinando el modelo generado con el gene y el modelo generado con el protein, extraído de un ejemplo de clase. En este modelo, la precisión aumenta, aunque la pérdida también lo hace. De todos modos, visto desde un punto de vista estadístico, no vemos una diferencia significativa ya que la precisión que hemos obtenido quedaría dentro del intervalo de confianza de la precisión del primero.

Como reflexión final, hemos visto que los autoencoders nos permiten reducir muchísimo el número de parámetros con los que la red neuronal densa va a trabajar. Al ejecutar el autoencoder, calculamos los pesos y después los congelamos, por lo que a partir de ahí, podemos conseguir trabajar en una dimensionalidad muchísimo más baja y, como hemos visto, obteniendo valores de precisión relativamente altos (cerca del 90%).